MARKETPLACE IMPLEMENTATION REPORT

Recommendation Algorithm Implementation for a Modern E-commerce Marketplace

A client-facing implementation report describing the recommendation and ranking system that was built, including smart scoring, anonymous-session behavior, impression tracking, aggregate metrics, and debugging support.

PREPARED BY
Arinze Obieze

3-Stage System
Candidate generation, scoring, and re-ranking

Cold-Start Ready
Supports first-time and anonymous users from day one

Production Ready
Implemented across storefront, analytics, and admin workflows

1. Executive Summary

The recommendation algorithm implemented for this marketplace is not a single “sort by newest” or “sort by popularity” rule. It is a structured recommendation system that retrieves relevant products, scores them across several commercial and trust dimensions, and then re-ranks the final list to improve list quality.

In practice, this implementation optimizes for more than one objective at the same time: relevance, purchase likelihood, product quality, seller trust, inventory health, freshness, diversity, and long-term marketplace health.

Implemented approach:

Built a three-stage recommendation architecture with contextual ranking for new users, anonymous session personalization for logged-out users, and stronger personalization for signed-in repeat users when behavioral history is available.

Relevance **Conversion** **Trust** **Diversity** **Exploration** **Fair Exposure**

2. Why The Marketplace Needed More Than Sorting

Basic sorting methods such as *newest first*, *price low to high*, and *highest rated* remain useful as customer controls, but they were too limited to function as the main ranking strategy of the marketplace.

A marketplace has to balance several realities at once:

- Some products are new and have no engagement history.
- Some sellers are operationally stronger and more reliable than others.
- Some products attract clicks but convert poorly.
- Some product lists look repetitive even if each item scores well individually.
- Different surfaces have different intent: home, category, search, store page, PDP recommendations, and account surfaces.

> The implemented recommendation system ranks by **score**, not by a single field in the database. The resulting marketplace ranking flow behaves as a decision engine, not just an ordering rule.

3. Implemented Architecture

STAGE 1

Candidate Generation

Retrieves a broad set of relevant products from search, category, popularity, store, and similarity pools before final ordering.

- Optimizes recall
- Fast and broad
- Does not decide final order

STAGE 2

Scoring

Assigns a score to each candidate using multiple signals such as relevance, quality, trust, conversion, stock, and freshness.

- Produces ranked preference
- Implemented as heuristic weighted scoring
- Structured for future model-based ranking

STAGE 3

Re-ranking

Improves the final list by enforcing diversity, seller caps, cold-start exposure, and commercial quality controls.

- Improves list quality
- Protects marketplace health
- Supports fairness and exploration

4. What The Algorithm Optimizes For

The implemented algorithm optimizes for meaningful commerce outcomes, not just surface-level engagement.

PRIORITY	OUTCOME	WHY IT MATTERS
1	Purchase likelihood	Direct business value and strongest signal of marketplace usefulness.
2	Add-to-cart likelihood	Useful leading indicator where purchase data is sparse.
3	Relevance	Ensures the products shown actually match user intent and page context.
4	Trust and quality	Reduces returns, disappointments, and marketplace credibility issues.
5	Healthy catalog exposure	Prevents ranking collapse into only a few old winners or a few dominant stores.

> If the system optimizes mostly for clicks, it will likely over-promote curiosity-driven products and under-optimize for purchase quality, trust, and long-term seller health.

5. Core Scoring Components

Relevance

Measures how well a product matches the query, category, surface, or related-item context.

Conversion

Estimates likelihood that the product leads to add-to-cart or purchase.

Popularity

Uses demand signals such as views, clicks, carts, wishlists, and recent purchases.

Product Quality

Captures rating, review volume, image completeness, content quality, and return risk.

Seller Trust

Reflects fulfillment quality, cancellation rate, dispute history, and seller verification.

Inventory Health

Protects the experience by favoring items with stable stock and low oversell risk.

Freshness

Allows newer products to enter the market without letting recency dominate forever.

Personalization / Business Signals

Uses user affinity where available, plus controlled strategic boosts when needed.

```
Base Score =
0.30 * Relevance
+ 0.20 * Conversion
+ 0.10 * Popularity
+ 0.10 * Product Quality
+ 0.10 * Seller Trust
+ 0.08 * Freshness
+ 0.07 * Inventory Health
+ 0.05 * Personalization

Served Score =
Base Score
+ Exploration Bonus
+ Merchandising Boost
- Return Risk Penalty
- Overexposure Penalty
```

The precise weights are being used as an explainable starting policy and can be tuned further using online experiments and observed business outcomes.

6. Different Surfaces Use Different Logic

SURFACE	PRIMARY GOAL	IMPLEMENTED RANKING EMPHASIS
Home page	Inspire and orient	Popularity, freshness, quality, diversity, strategic merchandising
Category page	Efficient browsing	Category relevance, conversion, stock health, trust, some freshness
Search results	Answer user intent	Strong relevance first, then conversion and trust
Product detail recommendations	Continue shopping	Similarity, compatibility, quality, complementarity
Store page	Represent seller well	Store bestsellers, recent arrivals, trusted items, good stock

7. Re-ranking Rules

After scoring, the system improves the list as a whole. This is the step that makes the product feed look intentional and premium instead of repetitive.

- Cap repeated exposure from the same seller in the first visible slots.
- Avoid near-duplicate products appearing side by side.
- Mix price bands so the feed does not become too expensive or too cheap.
- Mix subcategories or styles on broad browse pages.
- Reserve some space for newer, underexposed, but promising items.
- Downrank risky products and sellers with poor operational performance.

> Individual product strength is not enough. The implemented ranking system also produces a **good list**.

8. How It Works For The Very First User

A first-time user is a classic cold-start case. The system has little or no identity-based history, but it still has contextual information.

Signals still available for a first-time visitor:

- Landing page or entry surface
- Current category or collection
- Search query
- Device type
- Location or market context where relevant
- Traffic source or campaign source
- Immediate in-session actions

Implemented ranking mode:

The system uses contextual cold-start ranking. In other words, it acts like a smart merchandiser before acting like a personal assistant.

```
First-time user score =
0.30 relevance
+ 0.20 conversion
+ 0.15 popularity
+ 0.10 quality
+ 0.10 seller trust
+ 0.10 freshness
+ 0.05 inventory health
```

Example

If a first-time visitor lands on a “men’s loafers” search or category page, the system prioritizes the best matching loafers with strong conversion history, trusted sellers, solid stock, and a healthy spread in style and price. It does not keep pushing generic site-wide bestsellers that do not match the user’s immediate context.

9. How It Works For A User Who Is Not Logged In

A user who is not logged in is not the same thing as a first-time user. They may still have rich anonymous behavior attached to their browser session or device.

Useful anonymous signals:

- Recently viewed products
- Recent searches
- Clicked categories
- Time spent on PDPs
- Add-to-cart behavior
- Wishlist behavior where locally stored
- Preferred price band
- Session recency and page sequence

Implemented ranking mode:

The system uses anonymous session personalization on top of the contextual cold-start baseline.

The implementation assigns a stable anonymous ID and session ID in client storage so the recommendation system can build short-term preference memory before sign-in.

Example

A shopper is not logged in, but within 10 minutes they search for “black sneakers,” open four black footwear products, ignore premium price points above their pattern, and add one discounted pair to cart. The system responds by boosting similar sneakers, similar price bands, similar colors, and adjacent related items. This is meaningful personalization even without an account login.

10. Cold-Start Handling For New Products

New products have little or no behavioral history, so a performance-only ranking system would suppress them. This is why the implemented recommendation system uses a hybrid strategy.

- Use product attributes and content quality to estimate promise.
- Apply a temporary freshness boost.
- Allow measured exploration exposure.
- Use seller trust as a confidence anchor for new listings.

> A new product from a verified seller with complete imagery, solid pricing, clear copy, and healthy inventory receives some exposure even before it has accumulated purchase history.

11. Exploration vs Exploitation

A good marketplace recommendation system cannot only exploit what is already known to work. The implemented strategy also explores promising but uncertain products.

Exploitation

Show products that are already known to perform well on relevance, conversion, quality, and trust.

Exploration

Deliberately test new or underexposed products that appear promising but do not yet have enough evidence.

A practical policy is to keep the majority of exposure stable and exploitative, while reserving a small share, often 5-10%, for controlled exploration among new or underexposed quality products.

12. Fairness Across Sellers

The recommendation system supports marketplace health, not just short-term winner concentration. If the system always promotes only the largest or oldest sellers, catalog depth weakens and smaller high-quality sellers may disengage.

- Track seller exposure over time.
- Cap repeated exposure from one seller in top slots.
- Allocate some qualified exposure to underexposed but high-quality sellers.
- Use trust thresholds so fairness does not override safety and quality.

13. Data and Event Requirements

The recommendation system improves only if it learns from what was shown, what was ignored, and what converted. For that reason, impression-level logging was implemented as a core part of the rollout.

EVENT / DATASET	PURPOSE
Impressions	Teaches the system which products were shown and not only which were clicked.
Clicks / PDP views	Captures engagement and browsing interest.
Add-to-cart / wishlist	Strong signals of commercial intent.
Purchase events	Primary signal of recommendation value.
Returns / cancellations	Helps downrank low-quality or risky products and sellers.
Seller quality metrics	Supports trust-aware marketplace ranking.
Anonymous and signed-in profiles	Enables session-level and long-term personalization.

14. Implementation Summary

1. **Smart ranking API:** Implemented weighted scoring in the products API using relevance, conversion, popularity, quality, seller trust, freshness, inventory, personalization, discount, exploration, and risk controls.
2. **Anonymous context support:** Added anonymous and session-aware request headers so logged-out users can still receive adaptive recommendations.
3. **Impression tracking:** Implemented client-side product impression tracking with duplicate suppression per session and per surface position.
4. **Aggregate metrics pipeline:** Added a recommendation metrics table and refresh function to pre-compute 30-day product performance signals from analytics events.
5. **Operational tooling:** Added an admin ranking debug page so score breakdowns, ranking strategy, persona detection, and metrics source can be inspected during testing.

> The delivered system uses an explainable smart heuristic foundation and leaves a clear path for later learning-to-rank upgrades without changing the core recommendation flow.

15. Example End-To-End Flows

Example A: Brand New Visitor On Home Page

The system retrieves broad high-quality candidates from active categories, strong stores, recent arrivals, and popular products. It scores them on quality, popularity, trust, freshness, and stock. Then it re-ranks to ensure diversity across stores, categories, and price points.

Example B: Logged-Out User Browsing Bags

The user has no account session, but the platform sees recent bag searches, PDP views, and mid-price clicks. The system adapts by increasing bag affinity, preferred price range, and style similarity while preserving diversity and trust.

Example C: Signed-In Repeat User

The platform now uses both session data and long-term history, such as purchased categories, favorite stores, repeated styles, and price tolerance. Personalization becomes stronger, but relevance, trust, and list quality still remain guardrails.

16. Research References

The implementation choices in this report were informed by established research and production ranking practice.

- Google Research. *Recommending What Video to Watch Next: A Multitask Ranking System.* https://research.google/pubs/recommending-what-video-to-watch-next-a-multitask-ranking-system/
- Google Research. *TF-Ranking: Scalable TensorFlow Library for Learning-to-Rank.* https://research.google/pubs/tf-ranking-scalable-tensorflow-library-for-learning-to-rank-2/
- IJCAI 2022. *Multi-Stage Recommender Systems: A Survey / Review of architecture and ranking stages.* https://www.ijcai.org/proceedings/2022/771
- Google Research. *Learning to Rank with Selection Bias in Personal Search.* https://research.google/pubs/learning-to-rank-with-selection-bias-in-personal-search/
- Google Research. *Expediting Exploration by Attribute-to-Feature Mapping for Cold-Start Recommendations.* https://research.google/pubs/expediting-exploration-by-attribute-to-feature-mapping-for-cold-start-recommendations/
- Microsoft Research. *A Contextual-Bandit Approach to Personalized News Article Recommendation.* https://www.microsoft.com/en-us/research/publication/a-contextual-bandit-approach-to-personalized-news-article-recommendation-3/?lang=ja
- Google Research. *Towards Unified Metrics for Accuracy and Diversity for Recommender Systems.* https://research.google/pubs/towards-unified-metrics-for-accuracy-and-diversity-for-recommender-systems/
- User Modeling and User-Adapted Interaction. *Fairness and Exposure in Recommendation Systems.* https://link.springer.com/article/10.1007/s11257-023-09364-z

Final implementation summary: the marketplace now uses a multi-stage, context-aware, score-based recommendation system with explainable heuristic ranking, anonymous-session awareness, diversity-aware re-ranking, analytics-backed performance signals, and admin tooling for inspection.